

SHIVRAJ DHAYTADAK

+91 7972476081 | Email : shivraj.25d@gmail.com | LinkedIn: Shivraj Dhaytadak | Github: Shivraj-Dhaytadak

PROFILE SUMMARY

Dynamic **AI Strategy Consultant** with **4+ years of experience in designing and deploying Gen AI solutions**. Proven track record in optimizing workflows and reducing costs. Skilled in **Context Engineering, RAG, and Agentic AI** to automate complex workflows and generate high-quality outcomes.

SKILLS

Languages & Databases: Python, SQL (PostgreSQL, MySQL), NoSQL (MongoDB)

Agentic AI & LLM Frameworks: LangChain, LangGraph, CrewAI, MCP, A2A, RAG, Context Engineering, Guardrails, LangFuse, Prompt Engineering

Model Serving & ML: FastAPI, vLLM, PyTorch, Hugging Face Transformers, TRL (LoRA/QLoRA fine-tuning), Pydantic, Scikit-learn

Cloud & MLOps: Azure (AI Foundry, Function App, Blob Storage, Container Apps, APIM, Cosmos DB, SQL/Postgres, Application Insights); AWS (EC2, S3, Bedrock, SageMaker, RDS, Lambda, ECR, ECS, Fargate, CloudWatch)

DevOps & Tools: Git, GitHub Actions, Jenkins, Docker, Jira, Pytest

EXPERIENCE

Agentic AI Consultant

Jan 2026 - Present

Allianz Services

Pune, IN

- Engineered a **multi-layer Guardrails agent** using Azure Content Safety to screen 100k+ annual claims, instantly **halting malicious payloads and eliminating prompt injection vulnerabilities** across claim processing pipelines.
- Engineered a **LangChain-based PII redaction middleware** integrated with Azure Content Safety, automated data anonymization prior to persistence, and **ensured 100% compliance with Australia corporate privacy standards** across 100k+ annual insurance claim logs.
- Architected a **declarative JSON workflow engine** covering **19 claim journeys to process 100k+ annual claims**, slashing **new claim journey onboarding deployment time from 3 weeks to 1 day** without underlying codebase modifications.
- Designed a **two-tier prompt architecture decoupling domain-general guidelines from claim-specific rules**, eliminating 100% of prompt duplication while cutting agent configuration latency for 19 distinct claim journeys.
- Implemented **comprehensive Pydantic classes to handle output validations and LangGraph state**, with validations included in each agent node at runtime, eliminating downstream type issues caused by LLM non-deterministic behavior.

Data Scientist (Gen AI)

May 2025 - Jan 2026

Yash Technologies

Pune, IN

- Engineered a **high-performance FastAPI endpoint** integrating a Text-to-SQL graph, enabling seamless natural language query translation to structured SQL, **accelerating data retrieval by 40%**.
- Architected a **custom hybrid memory** framework (short-term & long-term) for chat sessions, enhancing contextual retention and summarization, **boosting query accuracy by 35% and improving SQL success rate scores by 25%**.
- Achieved **80% reduction in manual SQL effort and 35% higher response precision**, streamlining analytics workflows and decision-making.
- Built a multi-agent workflow using **CrewAI & Google Gemini 2.5 models** to support code generation and review across multiple programming languages.

Senior Software Engineer (Gen AI)

August 2022 - May 2025

Persistent Systems

Pune, IN

- Implemented high-performance model-serving endpoints using **FastAPI**, reducing **response latency by 50%** & reliable delivery for AI/ML applications.
- Created and executed automation scripts to replace repetitive tasks, enhancing precision in **data retrieval by 35%** and **boosting productivity by 25%** for product development initiatives.
- Developed an internal RAG codebase assistant leveraging LangChain and vector embeddings, reducing developer onboarding documentation lookup time by 40% across engineering teams.
- Enhanced LLM inference with RAG by leveraging document chunking, embedding optimization, and similarity search, **improving response relevancy by 40%**.
- Optimized RAG workflow by fine-tuning embeddings and retrieval strategies, **reducing latency by 30% and enhancing answer relevance**.
- Architected a data curation framework to aggregate and process datasets from multiple sources, **enabling 40% faster data integration to fine-tune LLMs, and improving processing accuracy by 30%**.

EDUCATION

Bachelor’s of Engineering in Computer Engineering
PES Modern College of Engineering, Pune

Pune, IN
August 2018 - July 2022

CERTIFICATION

Azure Certified AI Fundamentals - Issued August 2025
AWS Cloud Practitioner - Amazon Web Services - Issued April 2024 - Expires April 2027

AWARDS

Allianz Services - Best Team: H1 2026
Allianz Services - Best Performer: Q2 (July) 2026
Allianz Services - Best Performer: Q1 (March) 2026
Yash Technologies - Valuable Individual Asset Award: a direct reflection of the work put into developing Agentic AI workflows.
Persistent Systems - High Five Individual Award: for building multiple frameworks across teams and improving results through CI/CD.
Persistent Systems - Bravo Team Award: guided and mentored team members, fostering a collaborative and productive environment.